

Network Flows

Suppose that we are given the network in top of Figure 23.1, where the numbers indicate capacities, that is, the amount of flow that can go through the edge in unit time. We wish to find the maximum amount of flow that can go through this network, from S to T .

This problem can also be reduced to linear programming. We have a nonnegative variable for each edge, representing the flow through this edge. These variables are denoted f_{SA}, f_{SB}, \dots . We have two kinds of constraints: capacity constraints such as $f_{SA} \leq 5$ (a total of 9 such constraints, one for each edge), and flow conservation constraints (one for each node except S and T), such as $f_{AD} + f_{BD} = f_{DC} + f_{DT}$ (a total of 4 such constraints). We wish to maximize $f_{SA} + f_{SB}$, the amount of flow that leaves S , subject to these constraints. It is easy to see that this linear program is equivalent to the max-flow problem. The simplex method would correctly solve it.

In the case of max-flow, it is very instructive to “simulate” the simplex method, to see what effect its various iterations would have on the given network. Simplex would start with the all-zero flow, and would try to improve it. How can it find a small improvement in the flow? Answer: it finds a path from S to T (say, by depth-first search), and moves flow along this path of total value equal to the *minimum* capacity of an edge on the path (it can obviously do no better). This is the first iteration of simplex (see Figure 23.1).

How would simplex continue? It would look for another path from S to T . Since this time we already partially (or totally) use some of the edges, we should do depth-first search on the edges that have some *residual capacity*, above and beyond the flow they already carry. Thus, the edge CT would be ignored, as if it were not there. The depth-first search would now find the path $S - A - D - T$, and augment the flow by two more units, as shown in Figure 23.1.

Next, simplex would again try to find a path from S to T . The path is now $S - A - B - D - T$ (the edges $C - T$ and $A - D$ are full and are therefore ignored), and we augment the flow as shown in the bottom of Figure 23.1.

Next simplex would again try to find a path. But since edges $A - D$, $C - T$, and $S - B$ are full, they must be ignored, and therefore depth-first search would fail to find a path, after marking the nodes S, A, C as reachable from S . *Simplex then returns the flow shown, of value 6, as maximum.*

How can we be sure that it is the maximum? Notice that these reachable nodes define a *cut* (a set of nodes




Figure 23.1: Max flow

containing S but not T), and the *capacity* of this cut (the sum of the capacities of the edges going out of this set) is 6, the same as the max-flow value. (It must be the same, since this flow passes through this cut.) The existence of this cut establishes that the flow is optimum!

There is a complication that we have swept under the rug so far: when we do depth-first search looking for a path, we use not only the edges that are not completely full, but we must also traverse *in the opposite direction* all edges that already have some non-zero flow. This would have the effect of canceling some flow; canceling may be necessary to achieve optimality, see Figure 23.2. In this figure the only way to augment the current flow is via the path $S - B - A - T$, which traverses the edge $A - B$ in the reverse direction (a legal traversal, since $A - B$ is carrying non-zero flow).




Figure 23.2: Flows may have to be canceled

In general, a path from the source to the sink along which we can increase the flow is called an *augmenting path*. We can look for an augmenting path by doing for example a depth first search along the *residual network*, which we now describe. For an edge (u, v) , let $c(u, v)$ be its capacity, and let $f(u, v)$ be the flow across the edge. Note that we adopt the following convention: if 4 units flow from u to v , then $f(u, v) = 4$, and $f(v, u) = -4$. That is, we interpret the fact that we could reverse the flow across an edge as being equivalent to a “negative flow”. Then the *residual capacity* of an edge (u, v) is just

$$c(u, v) - f(u, v).$$

The residual network has the same vertices as the original graph; the edges of the residual network consist of all

weighted edges with strictly positive residual capacity. The idea is then if we find a path from the source to the sink in the residual network, we have an augmenting path to increase the flow in the original network. As an exercise, you may want to consider the residual network at each step in Figure 23.1.

Suppose we look for a path in the residual network using depth first search. In the case where the capacities are integers, we will always be able to push an integral amount of flow along an augmenting path. Hence, if the maximum flow is f^* , the total time to find the maximum flow is $O(Ef^*)$, since we may have to do an $O(E)$ depth first search up to f^* times. This is not so great.

Note that we do not have to do a depth-first search to find an augmenting path in the residual network. In fact, using a breadth-first search each time yields an algorithm that provably runs in $O(VE^2)$ time, regardless of whether or not the capacities are integers. We will not prove this here. There are also other algorithms and approaches to the max-flow problem as well that improve on this running time.

To summarize: the max-flow problem can be easily reduced to linear programming and solved by simplex. But it is easier to understand what simplex would do by following its iterations directly on the network. It repeatedly finds a path from S to T along edges that are not yet full (have non-zero residual capacity), and also along any reverse edges with non-zero flow. If an $S - T$ path is found, we augment the flow along this path, and repeat. When a path cannot be found, the set of nodes reachable from S defines a cut of capacity equal to the max-flow. Thus, *the value of the maximum flow is always equal to the capacity of the minimum cut*. This is the important *max-flow min-cut theorem*. One direction (that $\text{max-flow} \leq \text{min-cut}$) is easy (think about it: *any* cut is larger than *any* flow); the other direction is proved by the algorithm just described.

Duality

As it turns out, the max-flow min-cut theorem is a special case of a more general phenomenon called *duality*. Basically, duality means that for each maximization problem there is a corresponding minimizations problem with the property that any feasible solution of the min problem is greater than or equal any feasible solution of the max problem. Furthermore, and more importantly, *they have the same optimum*.

Consider the network shown in Figure 23.3, and the corresponding max-flow problem. We know that it can be written as a linear program as follows:




Figure 23.3: A simple max-flow problem

$$\begin{array}{rcll}
 \max & f_{SA} & + f_{SB} & \\
 & f_{SA} & & \leq 3 \\
 & & f_{SB} & \leq 2 \\
 & & & f_{AB} \leq 1 \\
 & & & f_{AT} \leq 1 \\
 & & & f_{BT} \leq 3 \\
 & f_{SA} & - f_{AB} & - f_{AT} = 0 \\
 & & f_{SA} & + f_{AB} & - f_{BT} = 0 \\
 & & & & f \geq 0
 \end{array} \quad P$$

Consider now the following linear program:

$$\begin{array}{rcll}
 \min & 3y_{SA} & + 2y_{SB} & + y_{AB} & + y_{AT} & + 3y_{BT} \\
 & y_{SA} & & & & + u_A & \geq 1 \\
 & & y_{SB} & & & + u_B & \geq 1 \\
 & & & y_{AB} & & - u_A & + u_B \geq 0 \\
 & & & & y_{AT} & - u_A & \geq 0 \\
 & & & & & y_{BT} & - u_B \geq 0 \\
 & & & & & & y \geq 0
 \end{array} \quad D$$

This LP describes the min-cut problem! To see why, suppose that the u_A variable is meant to be 1 if A is in the cut with S , and 0 otherwise, and similarly for B (naturally, by the definition of a cut, S will always be with S in the cut, and T will never be with S). Each of the y variables is to be 1 if the corresponding edge contributes to the cut capacity, and 0 otherwise. Then the constraints make sure that these variables behave exactly as they should. For example, the second constraint states that *if A is not with S , then SA must be added to the cut*. The third one states

that if A is with S and B is not (this is the only case in which the sum $-u_A + u_B$ becomes -1), then AB must contribute to the cut. And so on. Although the y and u 's are free to take values larger than one, they will be “slammed” by the minimization down to 1 or 0.

Let us now make a remarkable observation: these two programs have strikingly symmetric, *dual*, structure. This structure is most easily seen by putting the linear programs in matrix form. The first program, which we call the primal (P), we write as:

max	1	1	0	0	0		
	1	0	0	0	0	\leq	3
	0	1	0	0	0	\leq	2
	0	0	1	0	0	\leq	1
	0	0	0	1	0	\leq	1
	0	0	0	0	1	\leq	3
	1	0	-1	-1	1	$=$	0
	0	1	1	0	-1	$=$	0
	\geq	\geq	\geq	\geq	\geq		

Here we have removed the actual variable names, and we have included an additional row at the bottom denoting that all the variables are non-negative. (An unrestricted variable is denoted by unr.)

The second program, which we call the dual (D), we write as:

min	3	2	1	1	3	0	0		
	1	0	0	0	0	1	0	\geq	1
	0	1	0	0	0	0	1	\geq	1
	0	0	1	0	0	-1	1	\geq	0
	0	0	0	1	0	-1	0	\geq	0
	0	0	0	0	1	0	-1	\geq	0
	\geq	\geq	\geq	\geq	\geq	unr	unr		

Each variable of P corresponds to a constraint of D , and vice-versa. Equality constraints correspond to unrestricted variables (the u 's), and inequality constraints to restricted variables. Minimization becomes maximization. The matrices are transpose of one another, and the roles of right-hand side and objective function are interchanged.

Such LPs are called *dual* to each other. It is mechanical, given an LP, to form its dual. Suppose we start with a maximization problem. Change all inequality constraints into \leq constraints, negating both sides of an equation if

necessary. Then

- transpose the coefficient matrix
- invert maximization to minimization
- interchange the roles of the right-hand side and the objective function
- introduce a nonnegative variable for each inequality, and an unrestricted one for each equality
- for each nonnegative variable introduce a \geq constraint, and for each unrestricted variable introduce an equality constraint.

If we start with a minimization problem, we instead begin by turning all inequality constraints into \geq constraints, we make the dual a maximization, and we change the last step so that each nonnegative variable corresponds to a \leq constraint. Note that it is easy to show from this description that the dual of the dual is the original primal problem!

By the max-flow min-cut theorem, the two LPs P and D above have the same optimum. *In fact, this is true for general dual LPs!* This is the *duality theorem*, which can be stated as follows (we shall not prove it; the best proof comes from the simplex algorithm, very much as the max-flow min-cut theorem comes from the max-flow algorithm):

If an LP has a bounded optimum, then so does its dual, and the two optimal values coincide.